


Low latency inference for Transformers on Cloud

Sheik Mohamed Imran
AI Solutions Specialist - Intel
25-26 March, 2022



In deep learning applications, inference accounts for up to 90% of compute cost.

AWS ([Source](#))



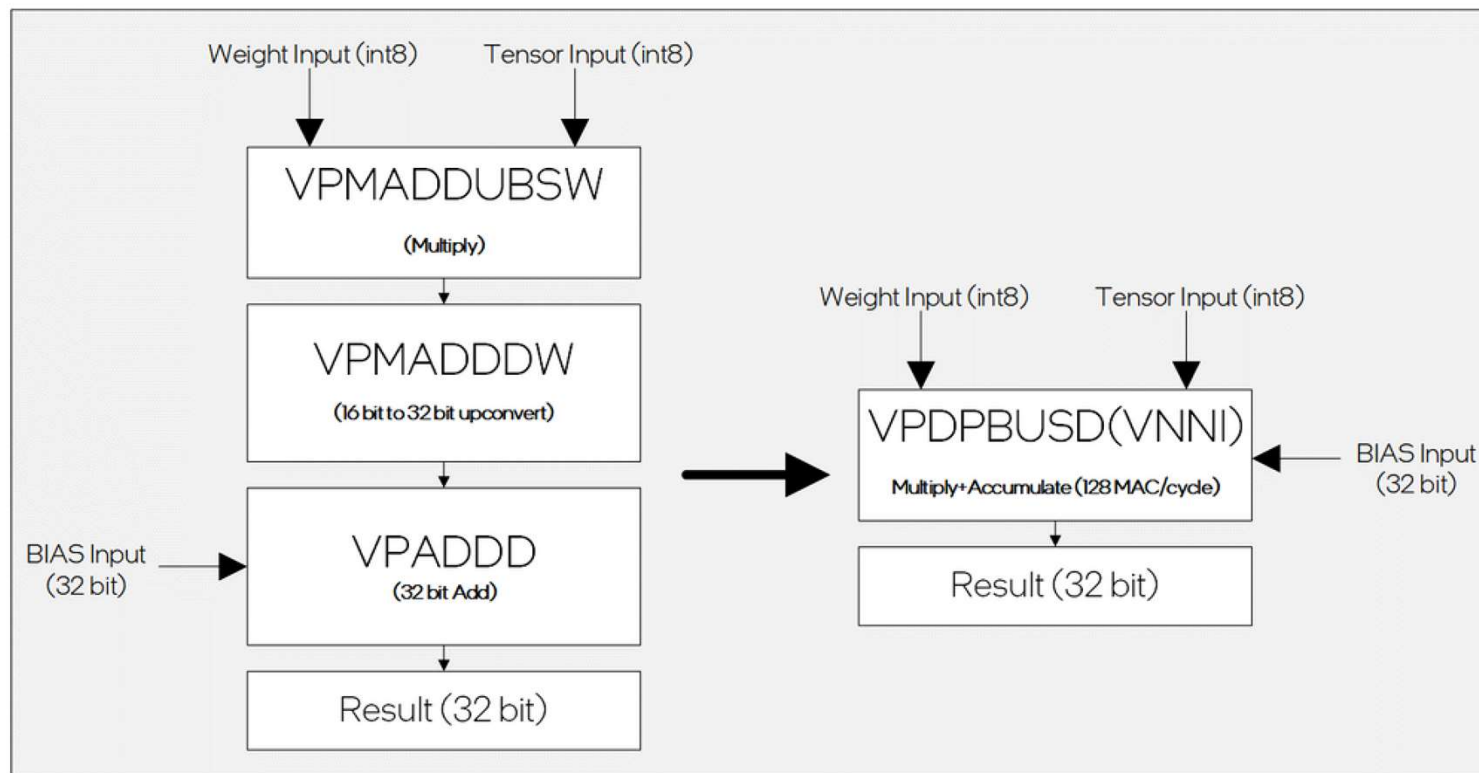
ONNX (Open Neural Network Exchange)

An open format to represent both deep learning and traditional models

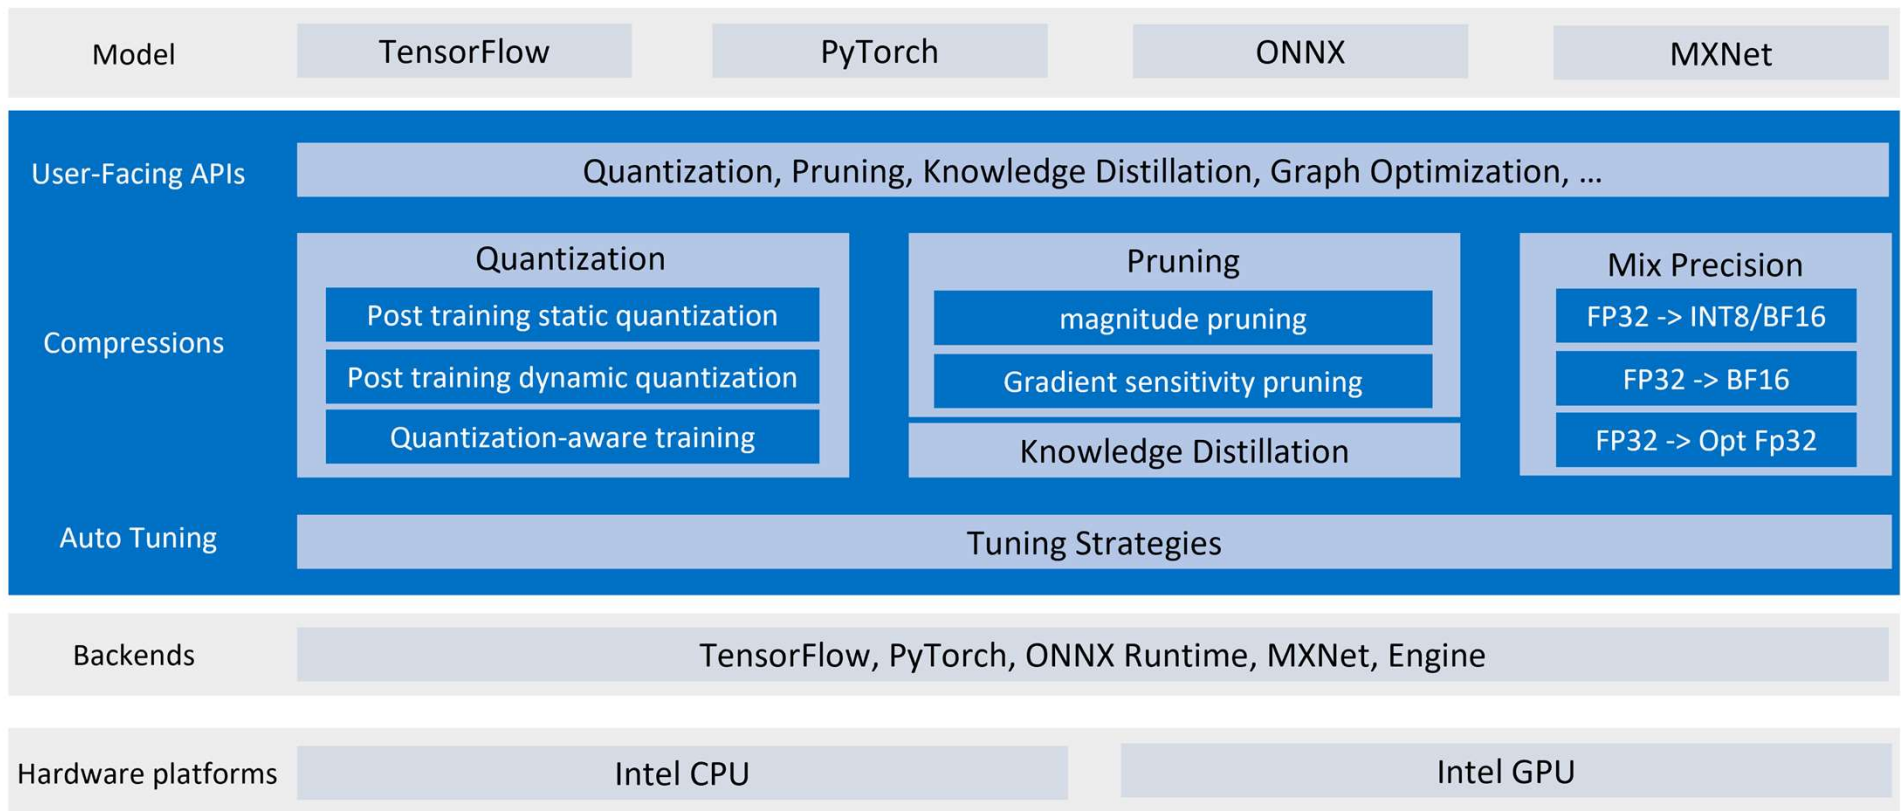
- Developed and maintained by community of partners (Microsoft, Facebook & AWS)
- Designed to offer interoperability across different frameworks.
- ONNXRuntime – Runtime library to maximize performance on Intel hardware for ONNX inference.



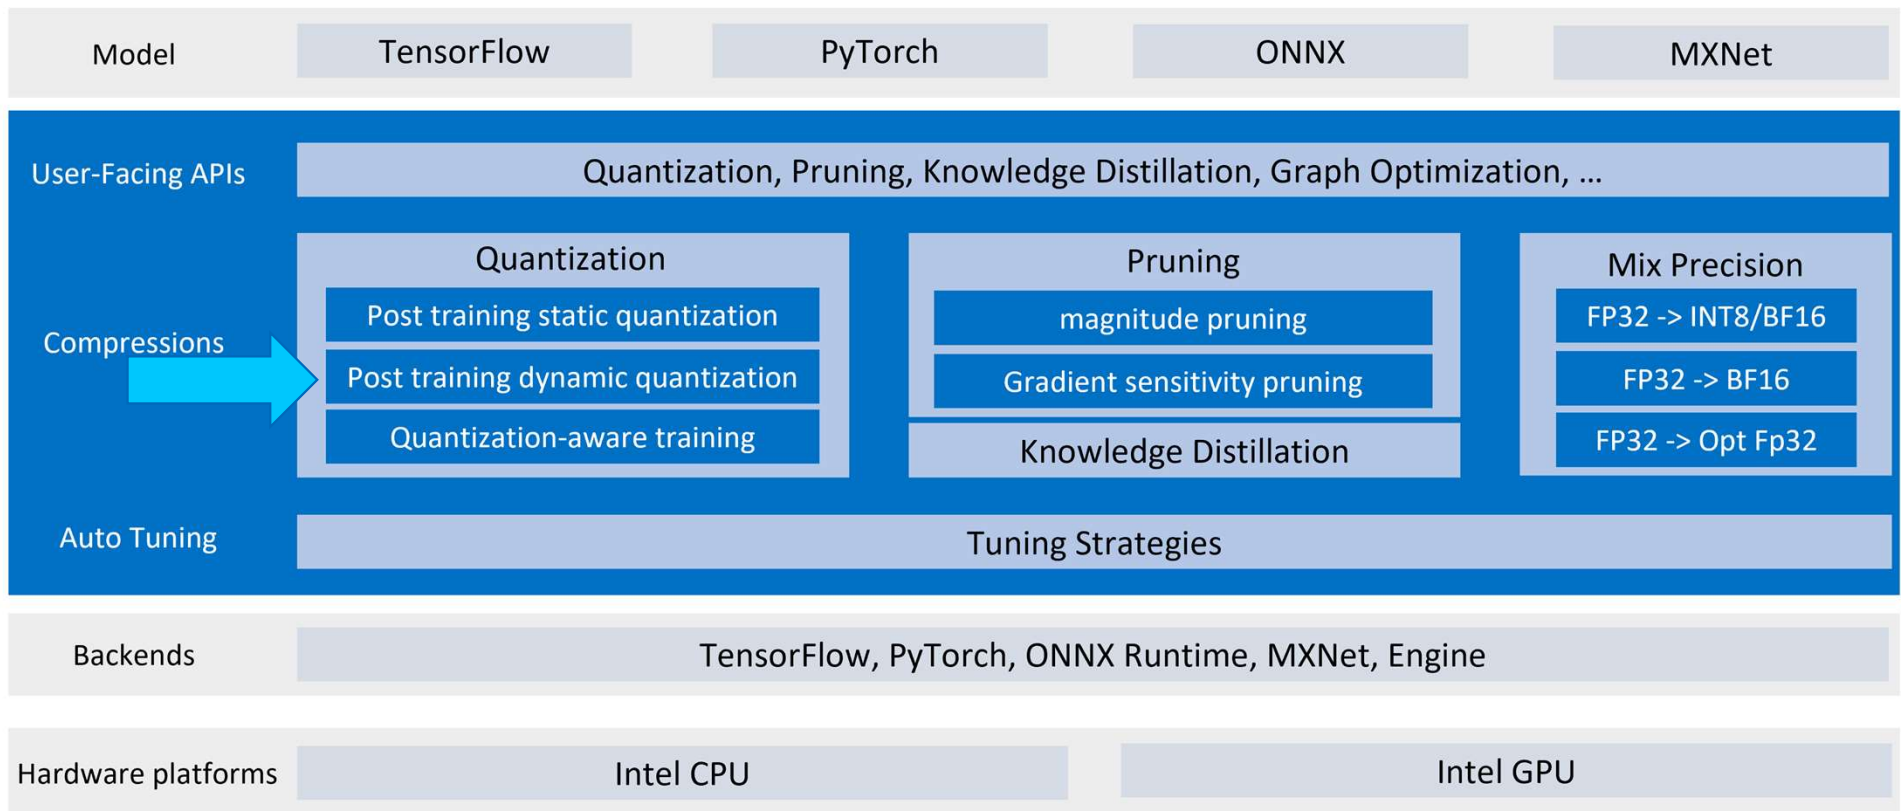
Intel® Deep Learning Boost - VNNI



Intel® Neural Compressor - Architecture

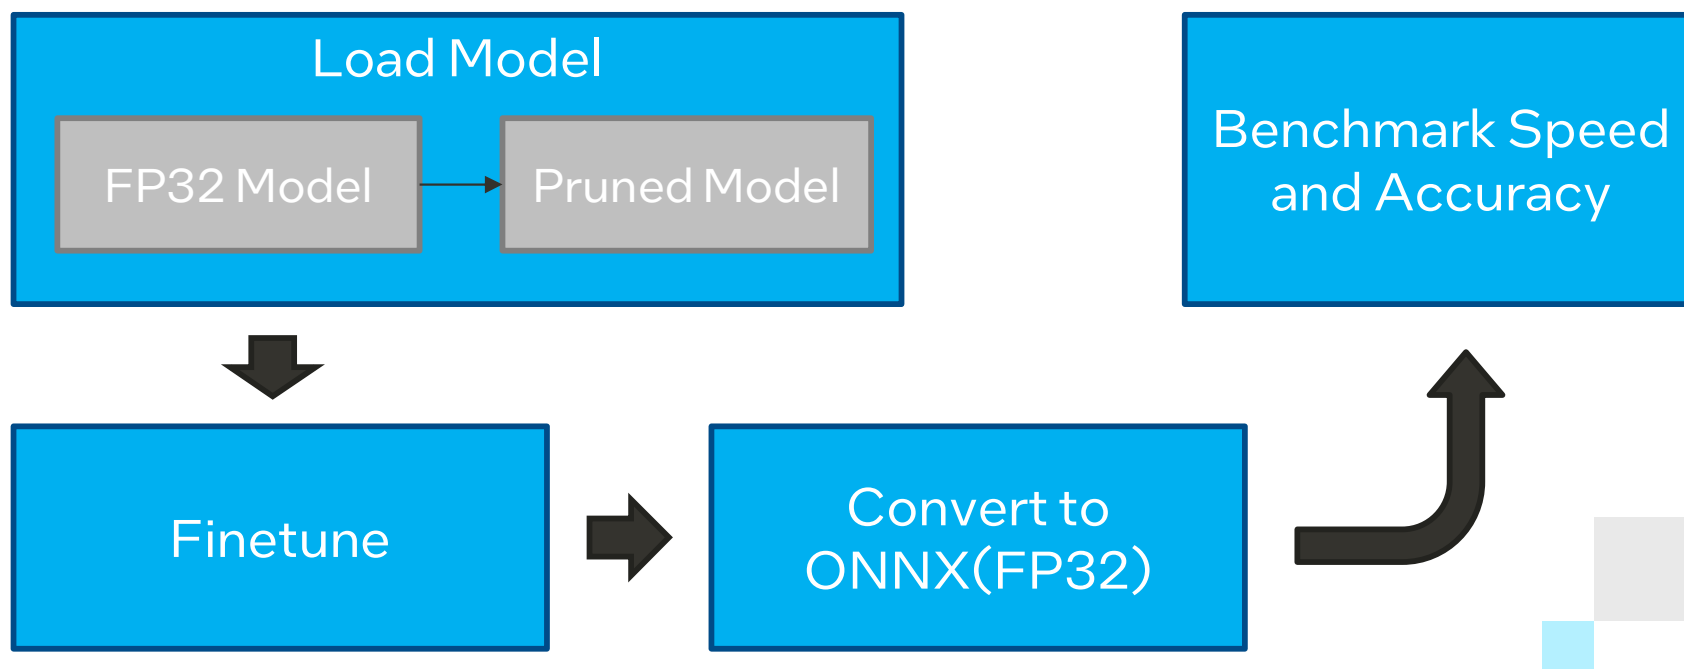


Intel® Neural Compressor - Architecture



Demo Overview


Part 1 – Finetune full & pruned models and run benchmarks.



Demo Overview

Model details:

- MiniLM – distilled approach created by Microsoft Research.
- The 12 Layer version from the paper shared in [microsoft/MiniLM-L12-H384-uncased](#).
- This is a 6-layer version of that model, by keeping only every second layer.



MINILM: Deep Self-Attention Distillation for Task-Agnostic Compression of Pre-Trained Transformers

Wenhui Wang Furu Wei Li Dong Hangbo Bao Nan Yang Ming Zhou
Microsoft Research
{wenwan, fuwei, lidong1, t-habao, nanya, mingzhou}@microsoft.com

nreimers / **MiniLM-L6-H384-uncased**

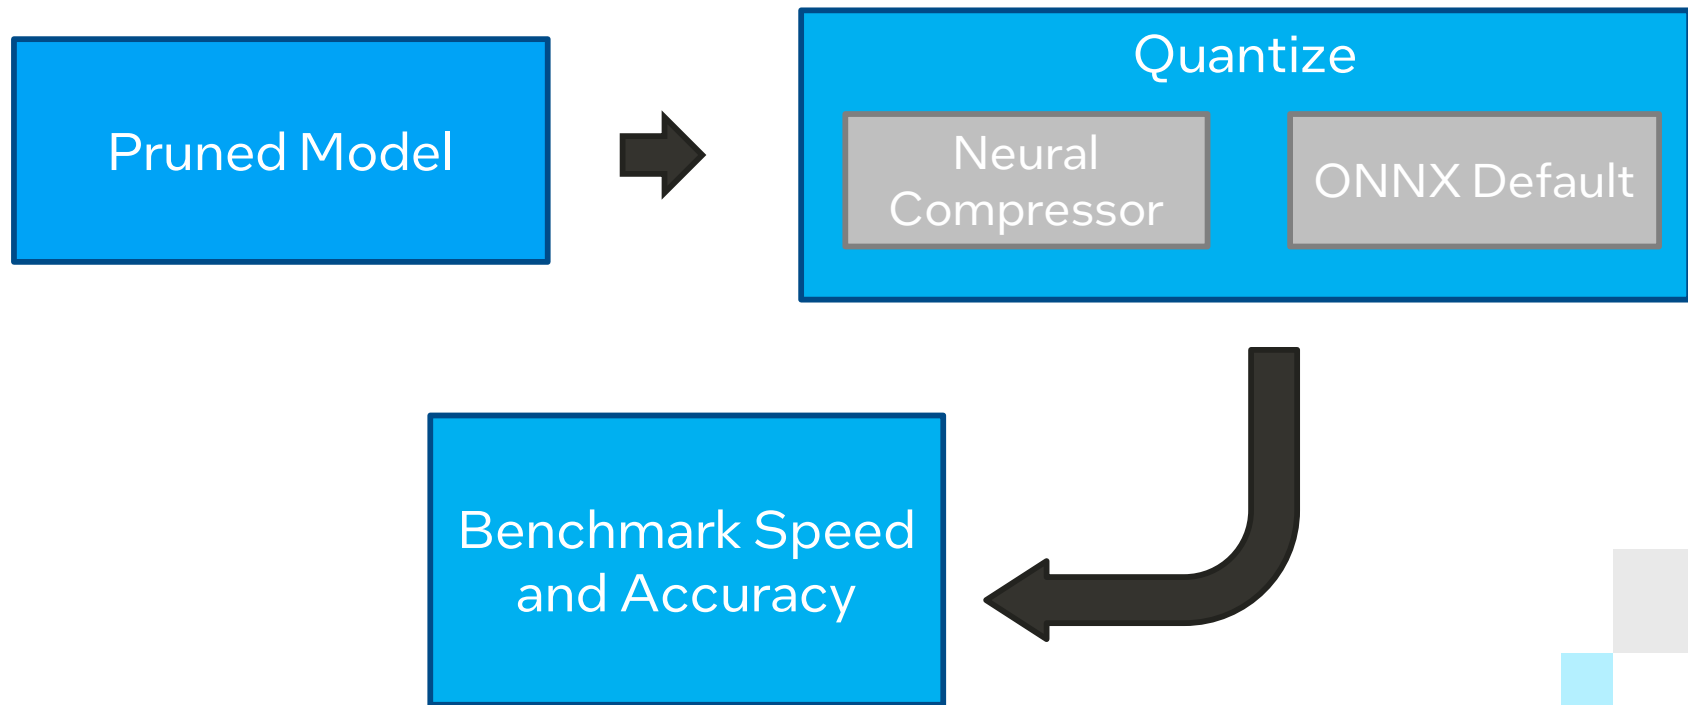
Feature Extraction PyTorch JAX Transfo

Model card Files and versions

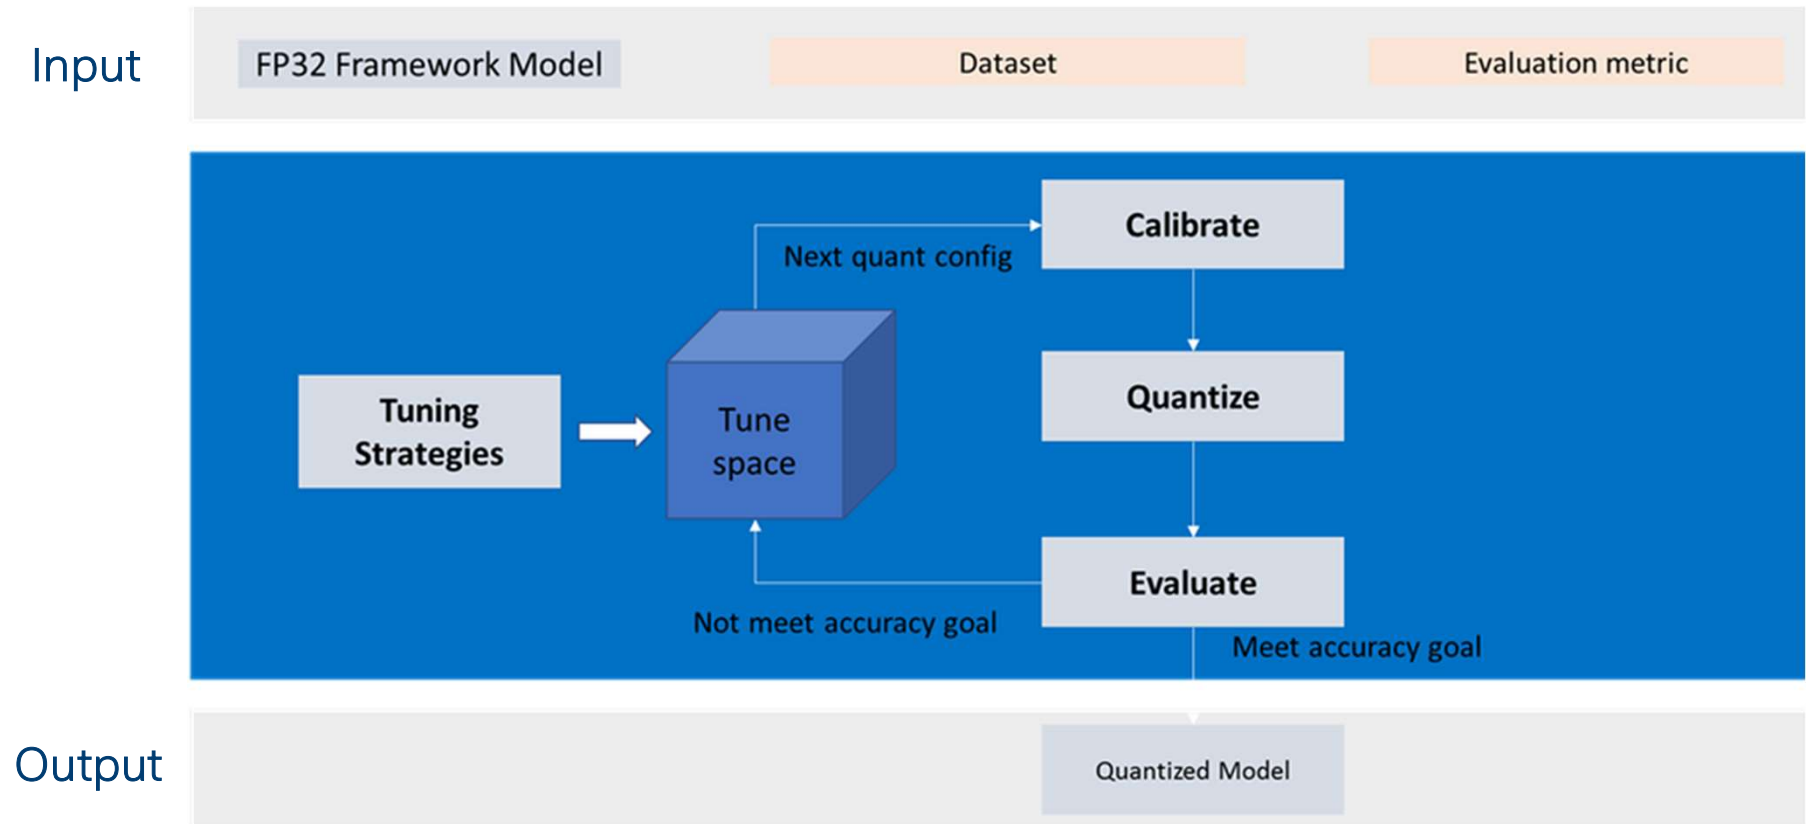
main MiniLM-L6-H384-uncased

Demo Overview

Part 2 – Covert pruned models to INT8 and run benchmarks.



Intel® Neural Compressor – Quantization Workflow



Intel® Neural Compressor - Configurations

Model Configurations

- “model” and “device” sections are mandatory
- “name” and “framework” fields are mandatory
- Possible values for framework are “tensorflow”, “mxnet”, “pytorch”, “pytorch_ipex”, “onnxrt_integerops” and “onnxrt_qlinearops”



```
model:
  name: bert
  framework: pytorch

device: cpu

quantization:
  approach: post_training_dynamic_quant

tuning:
  accuracy_criterion:
    relative: 0.01
  exit_policy:
    timeout: 0
    max_trials: 1200
  random_seed: 9527
```

Intel® Neural Compressor - Configurations

Quantization Configurations

- “approach” – to specify quantization method to be used post training (static/dynamic) quantization, or quantization aware training.



```
model:
  name: bert
  framework: pytorch

device: cpu

quantization:
  approach: post_training_dynamic_quant

tuning:
  accuracy_criterion:
    relative: 0.01
  exit_policy:
    timeout: 0
    max_trials: 1200
  random_seed: 9527
```

Intel® Neural Compressor - Configurations

Tuning Configurations

- “accuracy_criterion” could have “relative” or “obsolete” as fields.
- Example shown here allows relative accuracy loss of 1%
- “exit_policy” decides when to exit tuning. Here “timeout” (specified in seconds) is set at 0 which means early stop.
- “max_trials” indicates the maximum number of iterations to be tried.
- “random_seed” for deterministic tuning.



```
model:
  name: bert
  framework: pytorch

device: cpu

quantization:
  approach: post_training_dynamic_quant

tuning:
  accuracy_criterion:
    relative: 0.01
  exit_policy:
    timeout: 0
    max_trials: 1200
  random_seed: 9527
```



Further reading

- https://github.com/intel/neural-compressor/blob/master/docs/dynamic_quantization.md
- <https://community.intel.com/t5/Blogs/Tech-Innovation/Artificial-Intelligence-AI/Quantizing-ONNX-Models-using-Intel-Neural-Compressor/post/1355237>
- <https://pytorch.org/docs/stable/quantization.html>
- https://github.com/intel/neural-compressor/blob/master/docs/tuning_strategies.md

The background is a solid blue gradient. In the top-left corner, there is a cluster of four squares: a medium blue square, a light blue square, a dark blue square, and a small light blue square. In the bottom-right corner, there is another cluster of three squares: a light blue square, a medium blue square, and a dark blue square.

Thank you